

# DISTRIBUTED COMPUTING

Lecture I - Setting the Scene

Marco Massenzio

<http://codetrips.com>

<http://linkedin.com/in/mmassenzio>

# ROADMAP

- Why should we care?
- Principles of Distributed Computing
- MapReduce - Hadoop & friends
- Zookeeper - distributed computing patterns
- NoSQL - storing & retrieving data at scale
- Availability & Scalability - you really have to pick one
- CAP theorem - reality sucks
- Spark - streaming data for “quasi-realtime” big data analysis



# WHAT DO I NEED?

- Linux (Ubuntu) or Mac OSX laptop
- Oracle VirtualBox (free download)
- Python 3.4 (2.7 should work too)
- (some) Java (for MapReduce)
- (some) Familiarity with AWS (mostly EC2, S3, ELBs, Route53)
- Ability to download/install packages without much hand-holding
- (optional, recommended) Access to a good IDE (Eclipse or IDEA PyCharm/IntelliJ, free for students)
- (basic) Linux command-line (eg, ability to SSH into a remote instance, scp files across networks, mkdir, chown/chmod, very little else)
- Ability to use Google and StackOverflow



# WHY DO WE NEED DISTRIBUTED COMPUTING?

- Scalability
- Availability
- Performance



# SCALABILITY

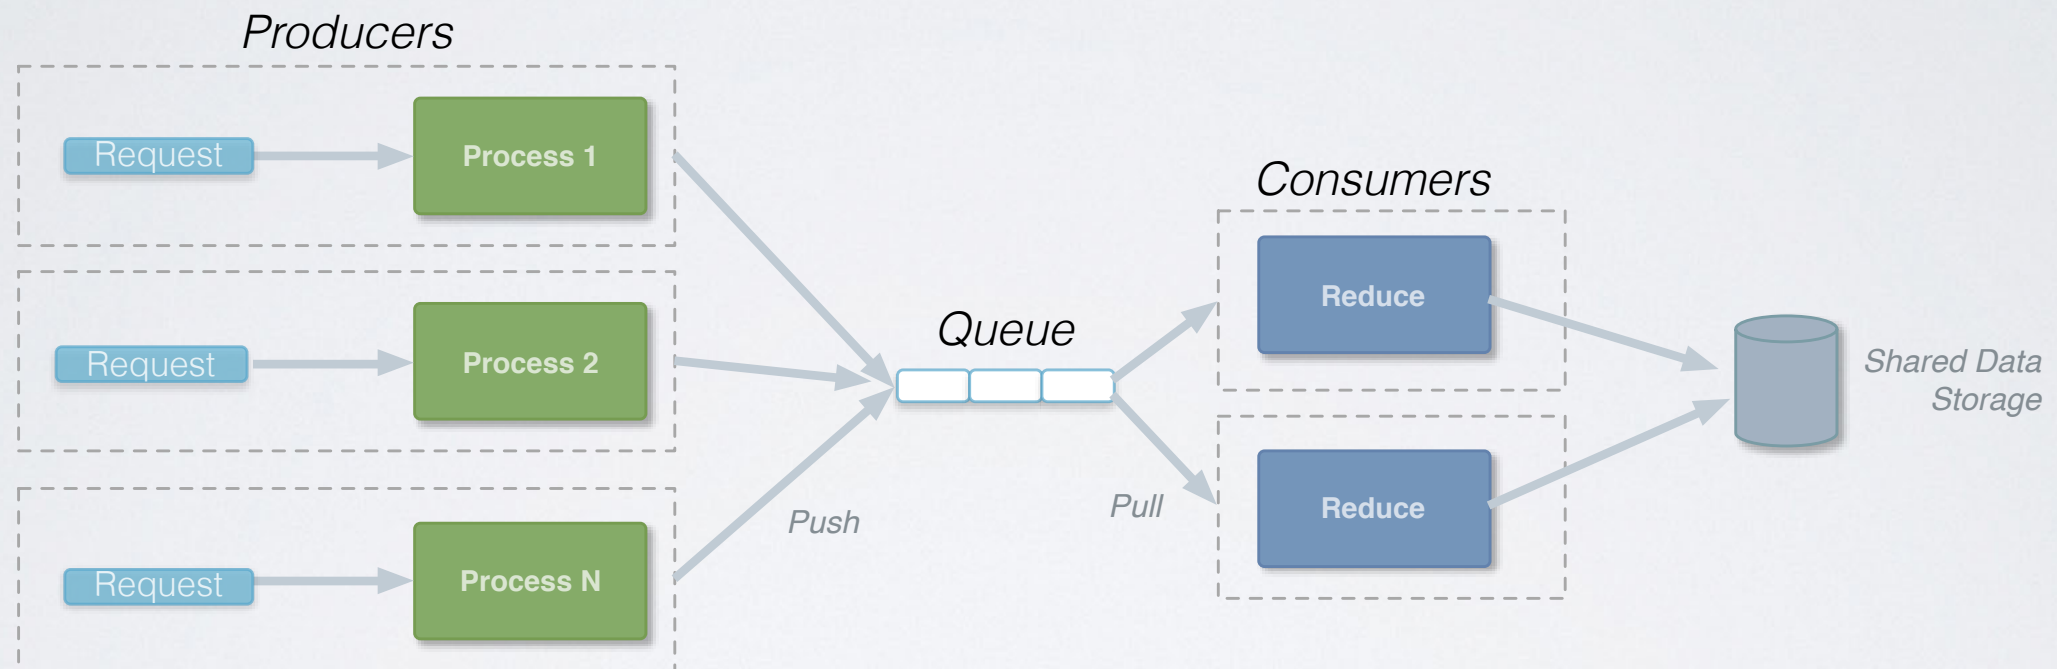
- Ability to serve a much higher rate of requests per seconds (typical metric: QPS) without significantly changing the architecture of the application
- Horizontal scalability: add more instances/processes of the same kind
- Vertical scalability: add more computing power and/or storage to each single instance
- They are not mutually incompatible; but require different architectural considerations (and usually are driven by very different requirements)
- **Beware:** adding more instances, one increases the probability that **any one of them** will fail (but increases the probability that **at least one** will not)



# MULTI-CORE CPUS

- For “CPU-bound” processes, adding more cores may make sense (provided the code can take advantage of that!)
- For I/O-bound processes this does not really help, unless one can distribute the data too (disk access is currently the limiting factor)
- Increased complexity of multi-threading code has led to consider different approaches (event-driven asynchronous architectures; actor-based systems; no-shared-state approaches)
- Distributed computing is usually taken to mean processes and systems based on multiple instances (physical or, more commonly, virtual) - it may relate to multi-process systems too

# CONSUMER/PRODUCER PATTERN



# SINGLE-PROCESS APPROACH

```
def get_n_samples(sensor, num):  
    """ Samples the cheap sensor and returns ``num`` values  
  
    :param sensor: the sensor to sample  
    :type sensor: Sensor  
    :param num: the number of samples to return, default 1,000  
    :return: the list of samples  
    :rtype: list of bool  
    """  
  
    count = num  
    samples = []  
    for x in sensor.get():  
        samples.append(x)  
        if count <= 0:  
            break  
        count -= 1  
    return samples
```

```
def should_run(sensors=3, samples=1000, faulty=1.0):  
    """ Finds out if we had a radioactive leak  
  
    We define a leak if more than half the sensors return an alarm, for more than three  
    consecutive samplings.  
  
    Naive implementation, samples the sensors and assumes they will all fit in memory.  
  
    :return: whether the sensor is faulty  
    :rtype: bool  
    """  
  
    num = sensors  
    sensors = [Sensor(faulty_pct=faulty) for _ in range(0, num)]  
    samples = [get_n_samples(s, num=samples) for s in sensors]  
    tot_count = len(samples[0])  
    for x in xrange(0, tot_count):  
        count = 0  
        for i in range(num):  
            if samples[i][x]:  
                count += 1  
        if count > 0:  
            logging.error("At sample %d, %s sensors were in the ALARM state", x, count)  
            if count > num / 2:  
                break  
    # Just because I wanted to show the use of for/else – a very Pythonic pattern!  
    else:  
        return False  
    return True
```



# MULTI-PROCESSING

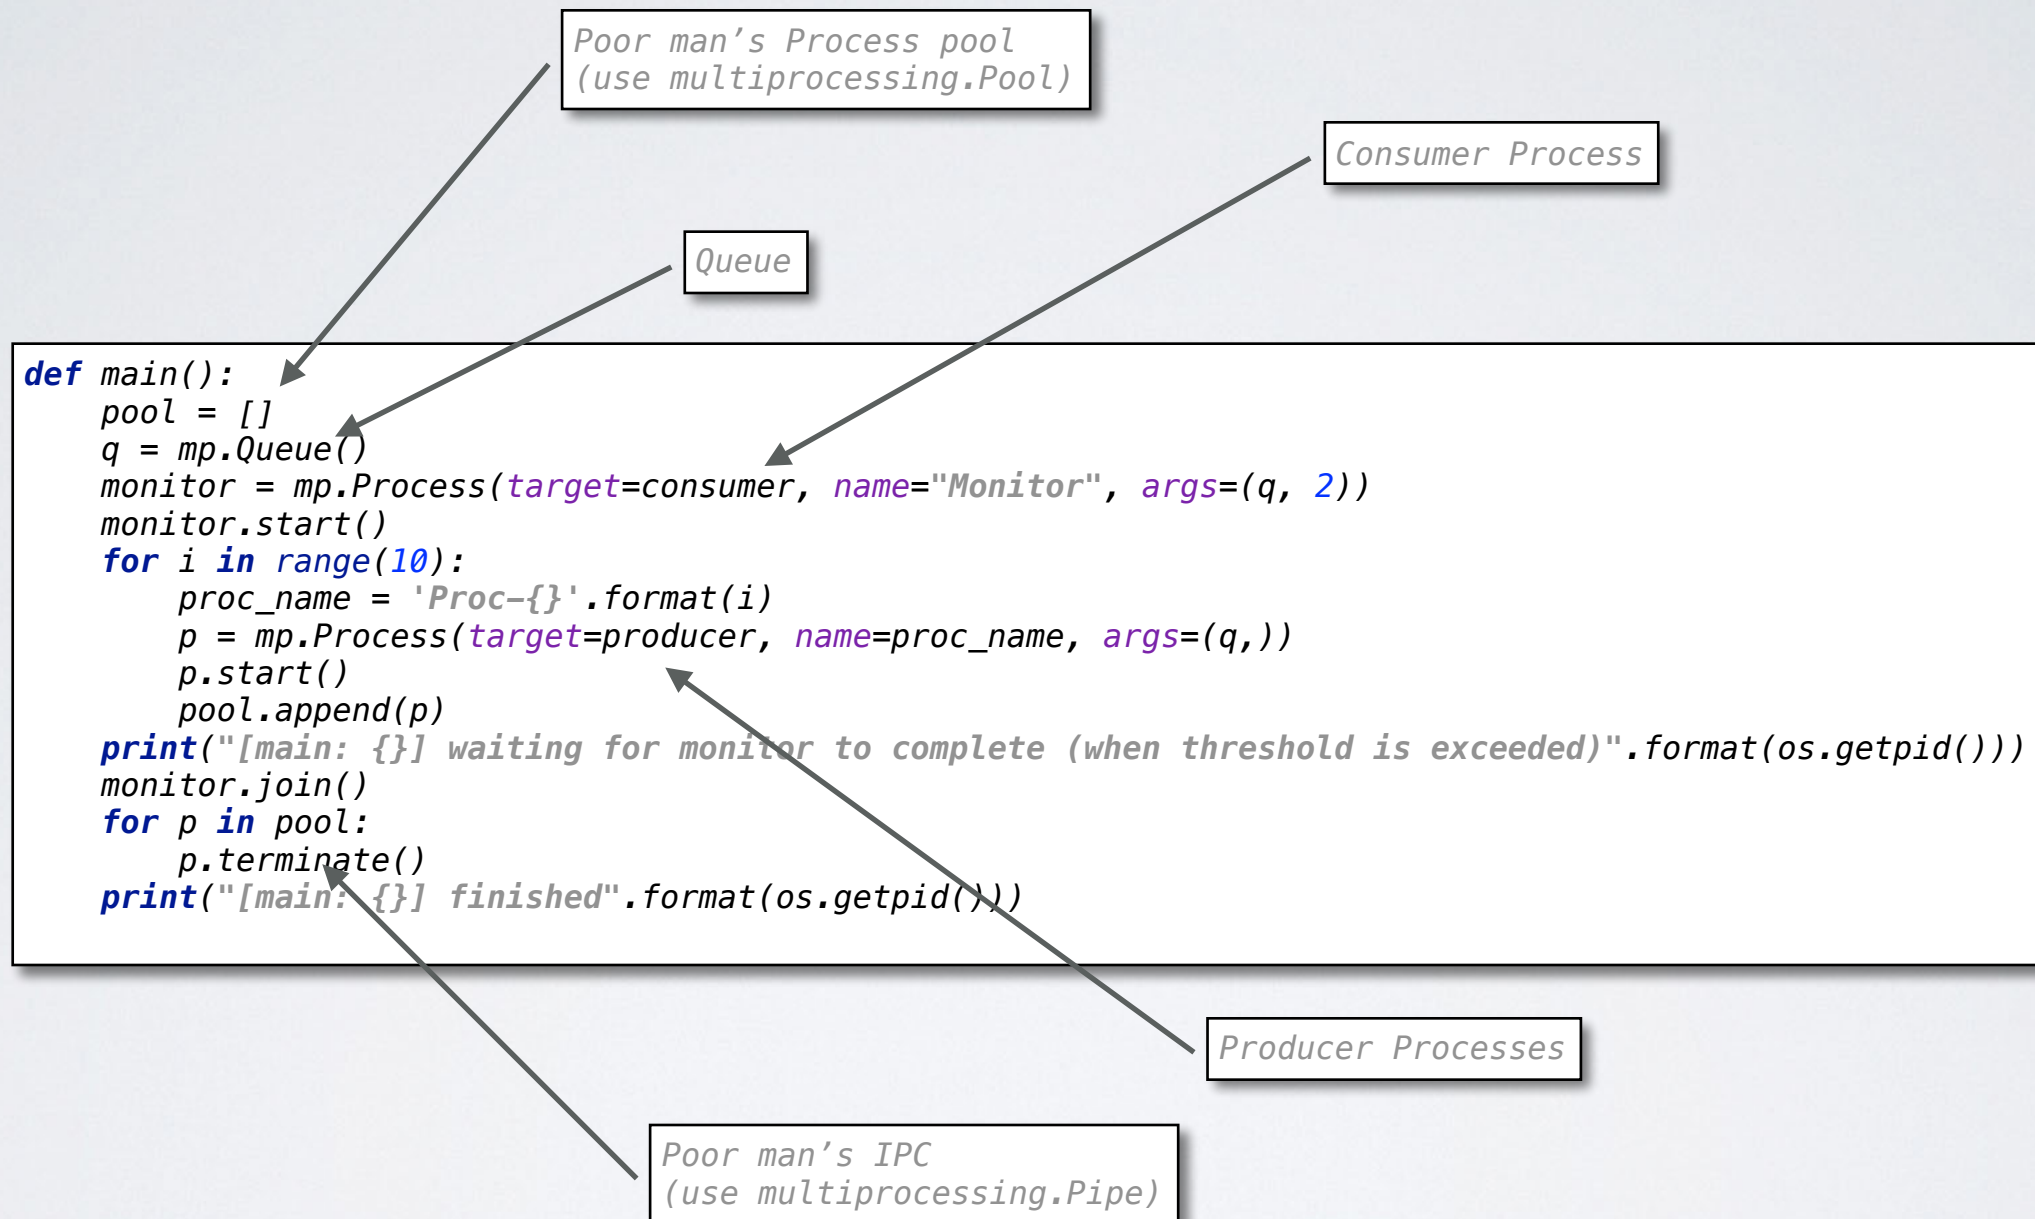
```
def producer(q, delay=0.500):
    """ It will forever put the sensor's readings onto the queue

    :param q: the queue to push sensor data to
    :param delay: between readings
    :return: None
    """
    print("[{}] producer started".format(os.getpid()))
    sensor = Sensor(faulty_pct=30.0)
    for x in sensor.get():
        q.put(x)
        time.sleep(delay)

def consumer(q, threshold=5):
    """ Reads values from the queue and raises an alarm if more than ``threshold``
        consecutive values are True

    :param q: the queue to read from
    :return: never, unless the threshold is exceeded
    """
    print("[monitor] Started with threshold {}".format(threshold))
    count = 0
    while count < threshold:
        reading = q.get(block=True)
        if reading:
            count += 1
        else:
            # reset the counter
            count = 0
    print("[monitor] Threshold exceeded - exiting")
```

# MULTI-PROCESSING (2)



# MULTIPROCESSING FOR REAL IS ACTUALLY (A LOT) MORE COMPLICATED

Shared state

```
def consumer(queue, idx, threshold=5, shared=None):
    """ Reads values from the queue and raises an alarm

    More than ``threshold`` consecutive values that are True will trigger an alarm.

    :param queue: the queue to read from
    :param threshold: The threshold at which we trigger the alarm, across ALL monitors
    :param shared: an optional shared ``Value`` for multiple Monitors
    :type shared: multiprocessing.Value
    :return: never, unless the threshold is exceeded
    """
    log("[monitor: {}] Started with threshold {}".format(os.getpid(), threshold))
    count = 0
    try:
        while shared.value < threshold:
            reading = queue.get(block=True)
            if reading:
                count += 1
                log('Alerting: {}'.format(count))
            else:
                # reset the counter
                count = 0
            if shared is not None:
                with shared.get_lock():
                    # NOTE: the double-check, as things may have changed between the test on the
                    # while and here; not doing this, causes some monitors to never terminate
                    if count == 0 and shared.value < threshold:
                        shared.value = 0
                    else:
                        shared.value += count
            log("[monitor-{}] Threshold exceeded - exiting".format(idx))
    except KeyboardInterrupt:
        # User pressed Ctrl-C, safe to ignore
        pass
```

MP Locking

Surprising facts



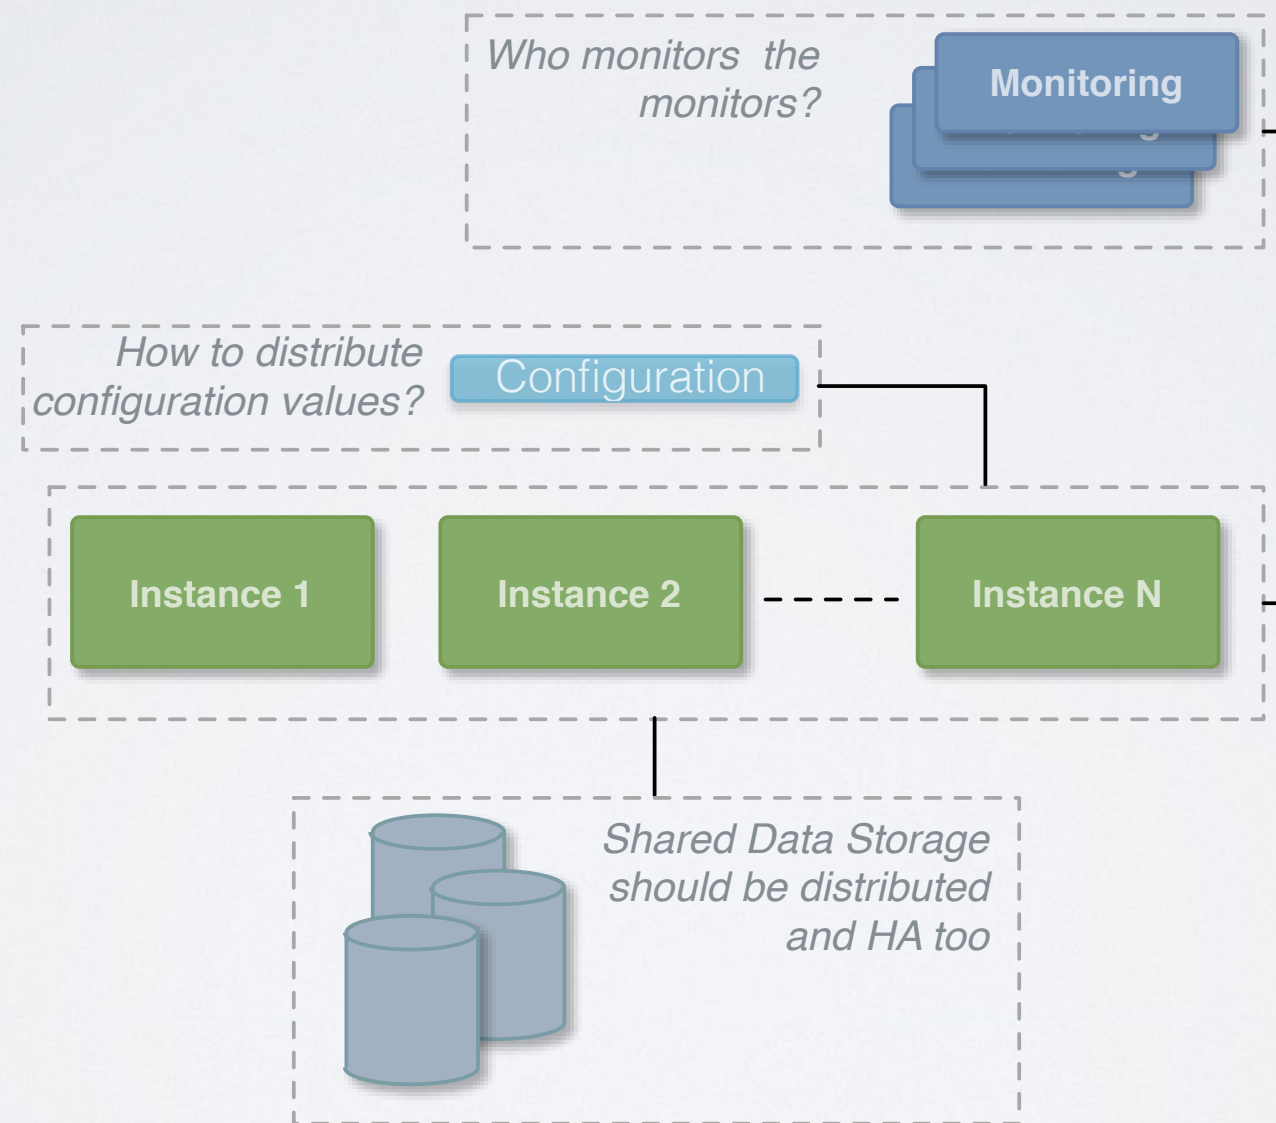
# DISTRIBUTED COMPUTING ACROSS MULTIPLE INSTANCES

For “horizontal” scalability and increased availability

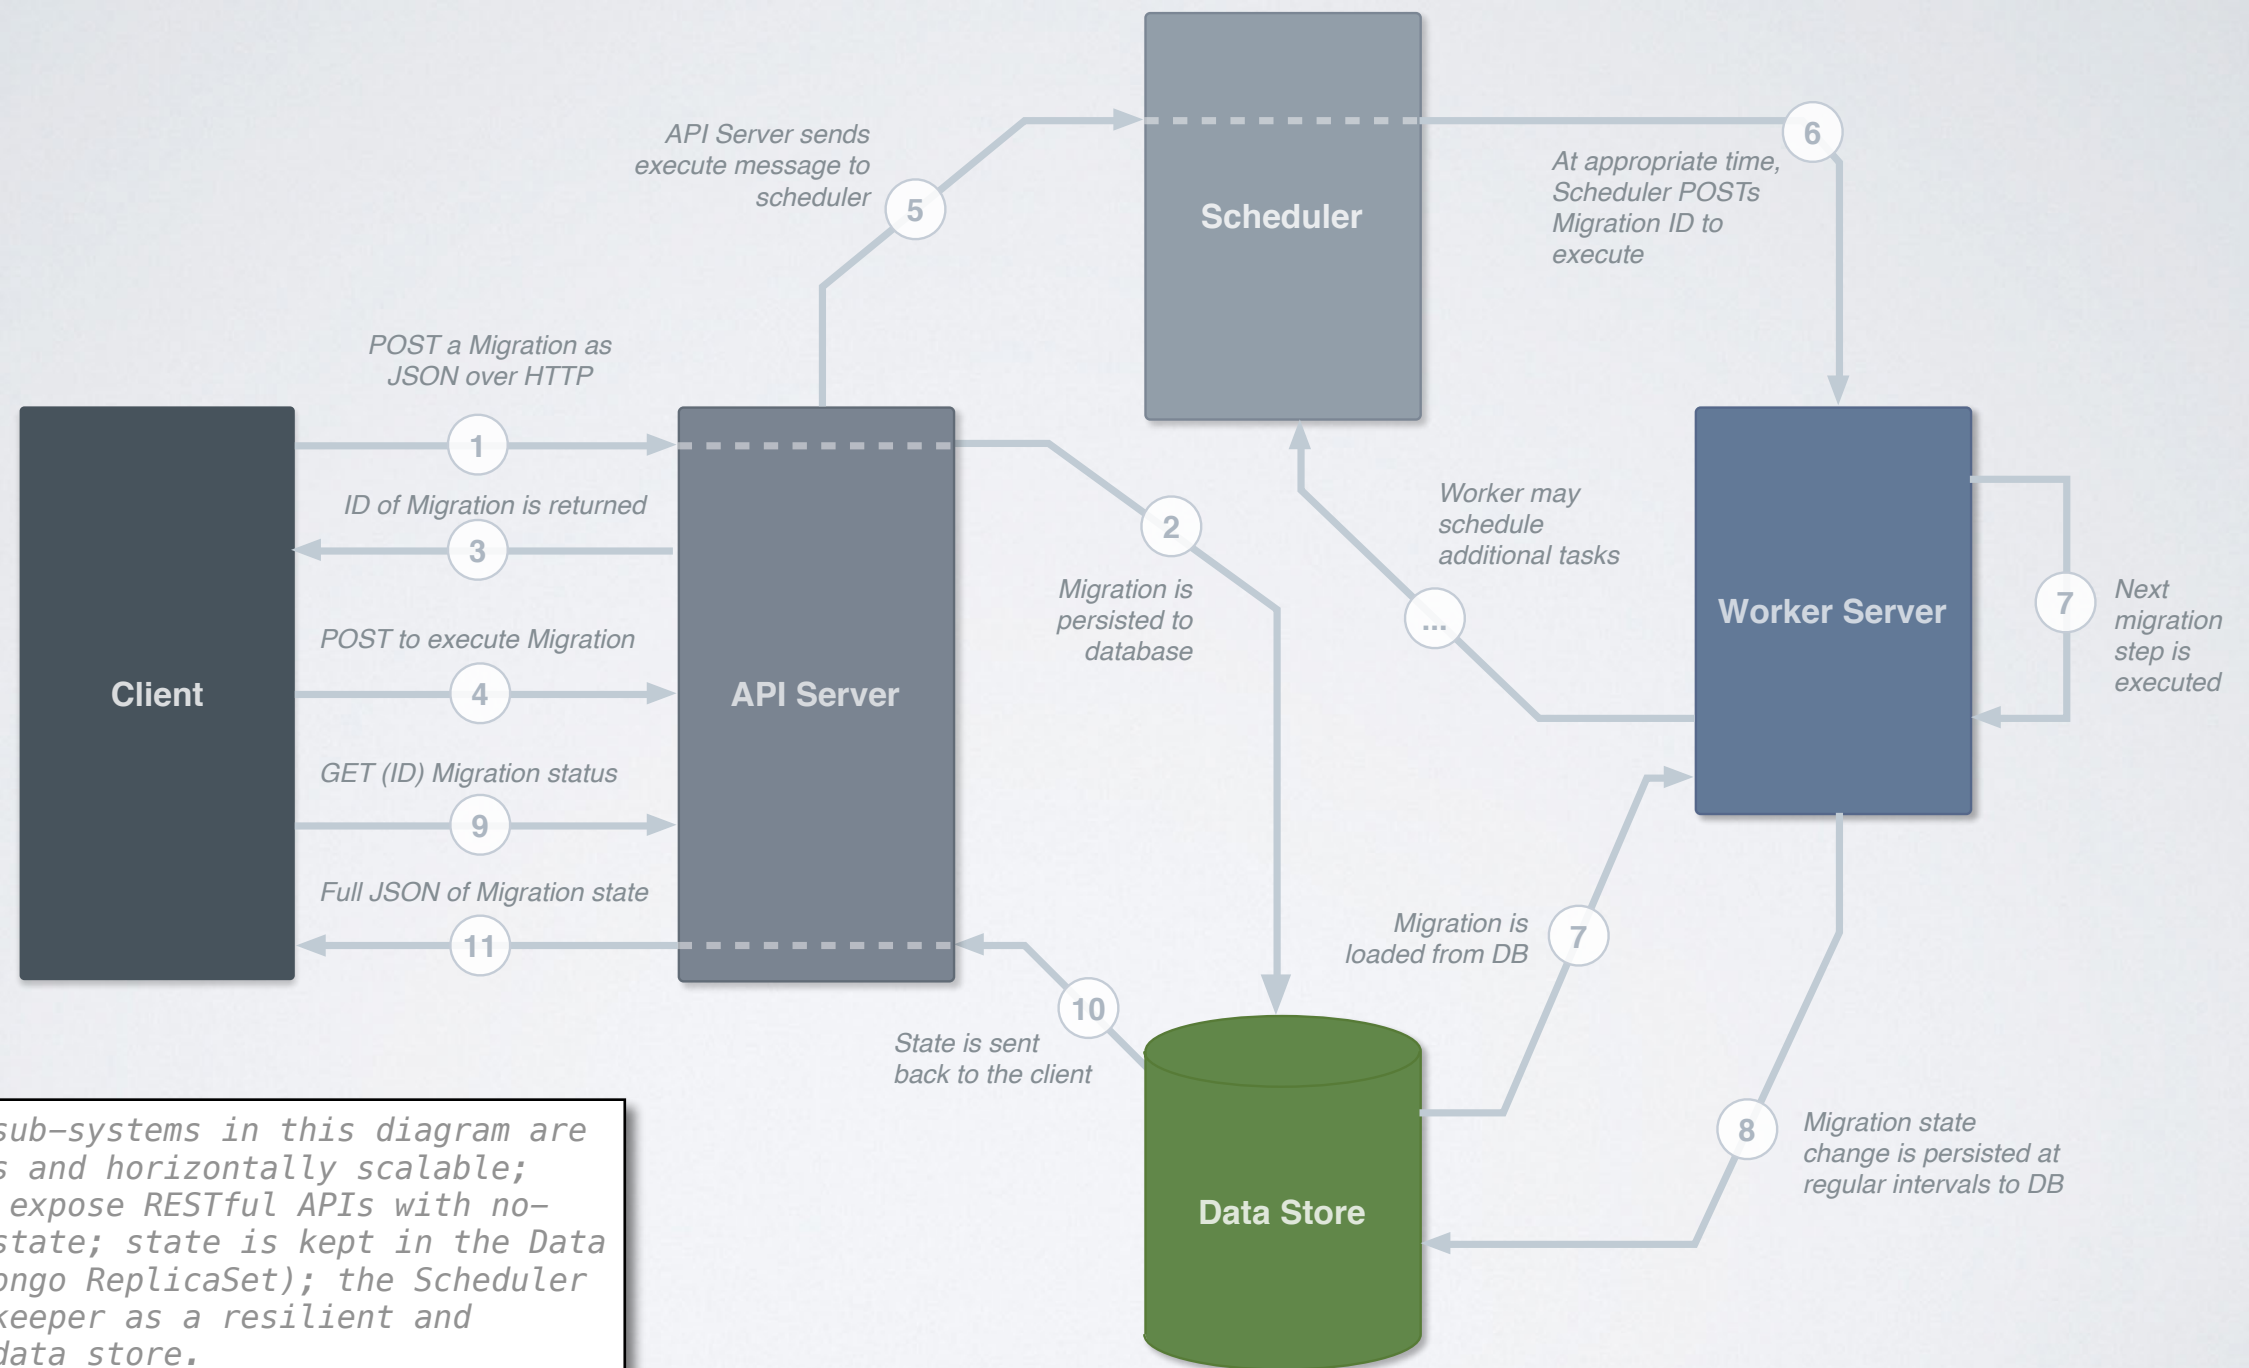
Challenges:

- Stateless architectures;
- Network protocol design (proprietary, or using well-established protocols: HTTP, ZMQ, TCP);
- Coordination & Monitoring of instances;
- Failure management & Restarting of instances
- Configuration management, resource discovery and fault diagnosis

# DISTRIBUTED COMPUTING ACROSS MULTIPLE INSTANCES (2)



# REAL-LIFE EXAMPLE: SEPARATION OF CONCERNS



All the sub-systems in this diagram are stateless and horizontally scalable; they all expose RESTful APIs with no-session state; state is kept in the Data Store (Mongo ReplicaSet); the Scheduler uses Zookeeper as a resilient and durable data store. Not shown, there are also a Monitoring subsystem and an Analytical Engine.



# SUMMARY

- Distributed computing is necessary to enable performance, availability and scalability of computing systems;
- Coordination, configuration and communication across distributed instances/ processes become major concerns;
- The computation model must change too - in particular, it is no longer safe to assume 'data locality' (or even its availability);
- Simplicity and homogeneity are two valuable attributes worth pursuing when designing distributed systems (think functional models);
- Synchronization and shared memory using multi-threading primitives fly in the face of "simplicity" and make the system more brittle (and more complicated - way more complicated - to diagnose and debug) - and may negatively impact on performance too.

# PROJECT

- Build various components over the span of the course
- Will require to interact with AWS APIs
- Build a “multiprocessing” computation
- Build a MapReduce (running against AWS EMR, using Python Streaming)
- Build a more complex MR, using a NoSQL DB (probably MongoDB) to store results
- Mostly meant to illustrate the issues to bear in mind when building a distributed system, rather than test your programming skills (but extra credit for clean, readable code!)
- All the sample code shown is on github: <https://github.com/massenz/MSAN694>
- Some of my rants: <http://codetrips.com>
- And some of my gists too: <https://gist.github.com/massenz>
- All of the slides will be posted on github too  
(and on the Course intranet, soon as I’m given access to it)